

# An Overview of Statistical Machine Learning Part II

Babak Shahbaba, Ph.D.

Associate Professor, Department of Statistics  
University of California, Irvine

Irvine, CA  
July 11, 2017

# Classification Models

# Logistic regression model

- When dealing with binary outcome variables, we assume the response variable,  $y_i$ , has a Bernoulli distribution (or Binomial if  $n_i > 1$ ),

$$y_i|\mu_i \sim \text{Bernoulli}(\mu_i)$$

- In this case, a common link function to connect the mean of the response variable to a set of predictors,  $x_i$ , is the *logit* function defined as follows:

$$g(\mu_i) = \log\left(\frac{\mu_i}{1 - \mu_i}\right) = \log\left[\frac{P(y_i = 1|x_i, \beta)}{1 - P(y_i = 1, \beta|x_i)}\right] = x_i\beta$$

where  $\beta = (\beta_0, \beta_1, \dots, \beta_p)$ .

- Note that

$$\mu_i = P(y_i = 1|x_i, \beta) = \frac{\exp(x_i\beta)}{1 + \exp(x_i\beta)}$$

# Interpretation

- To interpret  $\beta$ , notice that  $\log\left[\frac{P(y_i=1|x_i,\beta)}{1-P(y_i=1,\beta|x_i)}\right]$  is the log of odds for the outcome of interest,  $y_i = 1$ .
- The intercept  $\beta_0$  is therefore the log of odds when all predictors are set to zero (note that this might not make sense in some cases).
- Or we can say,  $\exp(\beta_0)$  is the odds when all predictors are set to zero.
- $\exp(\beta_j)$  on the other hand is how much the odds multiplicatively increases for one unit increase in  $x_j$  when all other predictors are fixed.
- Or we can say,  $\exp(\beta_j)$  is the odds ratio for subjects with  $X_j = x_j + 1$  compared to subjects with  $X_j = x_j$  when all other predictors are fixed.
- Positive  $\beta_j$  indicates that the odds increases as  $x_j$  increases (everything else fixed), where is for negative estimate of  $\beta_j$  the odds decreases as  $x_j$  increases (everything else fixed).

# Linear discriminant analysis

- When the set of  $p$  predictors,  $x$ , are continuous random variables, we can assume that their joint distribution is multivariate normal for each class,

$$f_k(x) = (2\pi)^{-p/2} |\Sigma|^{-1/2} \exp\left[-\frac{1}{2}(x - \mu_k)^T \Sigma^{-1}(x - \mu_k)\right]$$

- Note that in this setting, only the mean of the distributions,  $\mu_k$ , changes from one class to another. The covariance matrix  $\Sigma$  remains the same for all classes.
- This assumption is of course not realistic and is made only for simplicity. We will relax it later.

# Linear discriminant analysis

- Using Bayes theorem, we have

$$P(y = k|x) = \frac{\pi_k f_k(x)}{\sum_{k'=1}^K \pi_{k'} f_{k'}(x)}$$

where  $\pi_k = P(y = k)$ .

- For a given value of  $x$ , the denominator remains the same for all classes. Therefore, we can define the discriminant function based on the numerator,  $\pi_k f_k(x)$ , or more commonly based on its log,

$$\begin{aligned}\delta_k(x) &= \log \pi_k + \log[f_k(x)] \\ &= \log \pi_k - \frac{1}{2} \log(|\Sigma|) - \frac{1}{2} (x - \mu_k)^T \Sigma^{-1} (x - \mu_k)\end{aligned}$$

# Linear discriminant analysis

- With further simplification (and removing the constant parts), we have

$$\delta_k(x) = \log \pi_k - \frac{1}{2} \mu_k^T \Sigma^{-1} \mu_k + x^T \Sigma^{-1} \mu_k$$

- Note that the above functions are linear in  $x$ .
- Therefore, we refer to them as *linear discriminant functions*.
- Classifying cases according to these functions is called *linear discriminant analysis* (LDA).

# Linear discriminant analysis

- We can estimate  $\pi_k$  and  $\mu_k$  for  $k = 1, \dots, K$ , and  $\Sigma$  as follows:

$$\hat{\pi}_k = \frac{n_k}{n}$$

$$\hat{\mu}_k = \frac{1}{n_k} \sum_{i:y_i=k}^{n_k} x_i$$

$$\hat{\Sigma} = \frac{1}{n - K} \sum_{k=1}^K \sum_{i:y_i=k}^{n_k} (x_i - \hat{\mu}_k)(x_i - \hat{\mu}_k)^T$$

where  $n_k$  is the number of observed cases (training cases) that belong to class  $k$ .



# Linear discriminant analysis

- After estimating the model parameters, we assign each case,  $i$ , to the class whose value of the discriminant function,  $\delta_k(x_i)$ , is the highest.
- Cases for which  $\delta_k(x) = \delta_l(x)$  fall on the decision boundary between the two classes  $k$  and  $l$ .
- For these cases,  $\delta_k(x) - \delta_l(x) = 0$ , which means

$$\log \frac{\pi_k}{\pi_l} - \frac{1}{2}(\mu_k - \mu_l)^T \Sigma^{-1}(\mu_k - \mu_l) + x^T \Sigma^{-1}(\mu_k - \mu_l) = 0$$

- Note that the above equation, which specifies the decision boundary, is linear in  $x$ . As the result, the decision boundaries are *hyperplanes* in the  $p$ -dimensional space. (The decision boundary is straight line if we have two predictors only.)

# Linear discriminant analysis

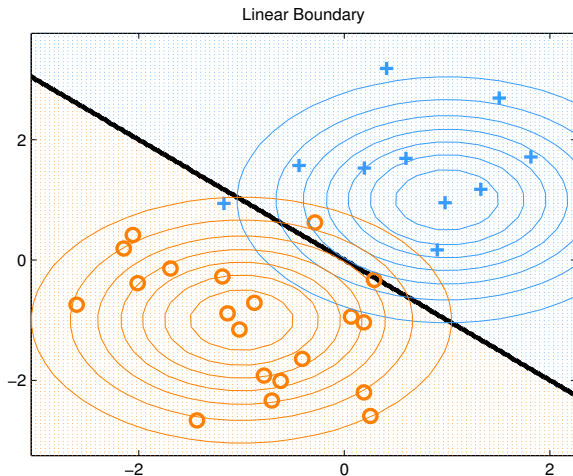


Figure: Figure 4.5a in Murphy (2012)

# Quadratic discriminant analysis

- As mentioned above, the equal-covariance assumption is restrictive and is only made for convenience.
- By relaxing this assumption, the discriminant function becomes

$$\delta_k(x) = \log \pi_k - \frac{1}{2} \log(|\Sigma_k|) - \frac{1}{2} (x - \mu_k)^T \Sigma_k^{-1} (x - \mu_k)$$

which are quadratic functions of  $x$ ; hence, they are called *quadratic discriminant functions*.

- Classifying cases according to these functions is called *quadratic discriminant analysis* (QDA).
- The decision boundaries for this approach are not linear any more.

# Linear discriminant analysis




Figure: Figure 4.3a in Murphy (2012)

# Naive Bayes models

- This is an alternative classification model, which is especially attractive when the dimension  $p$  is large.
- In this approach, we again use Bayes theorem to obtain the probability of each class given the observed values of predictors,

$$P(y = k|x_1, \dots, x_p) = \frac{P(y = k)P(x_1, \dots, x_p|y = k)}{\sum_{k'=1}^K P(y = k')P(x_1, \dots, x_p|y = k')}$$

- This time, however, we make an assumption that is naive and possibly wrong, but it simplifies the model: we assume that given a class  $y = k$ , the predictors are independent,

$$P(x_1, \dots, x_p|y = k) = \prod_{j=1}^p P(x_j|y = k)$$

# Naive Bayes models

- As a result of the above naive assumption, the model simplifies to

$$P(y = k|x_1, \dots, x_p) = \frac{P(y = k) \prod_{j=1}^p P(x_j|y = k)}{\sum_{k'=1}^K P(y = k') \prod_{j=1}^p P(x_j|y = k')}$$

- As before, we assign each case,  $i$ , to the class with the highest conditional probability given  $x_{i1}, \dots, x_{ip}$ .
- It is more common to distinguish between two classes using the following logit function

$$\begin{aligned} \log \frac{P(y = k|x_1, \dots, x_p)}{P(y = l|x_1, \dots, x_p)} &= \log \frac{P(y = k) \prod_{j=1}^p P(x_j|y = k)}{P(y = l) \prod_{j=1}^p P(x_j|y = l)} \\ &= \log \frac{\pi_k}{\pi_l} + \sum_{j=1}^p \log \frac{P(x_j|y = k)}{P(x_j|y = l)} \end{aligned}$$

# Naive Bayes models

- In practice, we estimate  $\pi_k$  using the proportion of observed cases that belong to class  $k$ .
- To estimate  $P(x_j|k)$ , we first need to assume a probability distribution model for  $x_j$  given  $k$ .
- If  $x_j$  is categorical, we can estimate  $P(x_j|k)$  using the observed proportion of each category of  $x_j$  for cases with  $y = k$ .
- If  $x_j$  is continuous, we can assume  $x_j|k$  has a Gaussian distribution and estimate its mean and variance using the cases with  $y = k$ .

# Tree Models



# Tree Models

- Decision trees are models that recursively partition the input space into regions and define a local map between each region and the response variable.





Figure: Figure 16.1 in Murphy (2012).

- Starting from the root, when a rule is satisfied, we move to the left branch; otherwise, we move to the right branch.
- When we reach a leaf, we use the subset of data, which fall in the corresponding region, to estimate the response variable.
- The corresponding model can be presented as follows:

$$\hat{y}_i = \sum_{m=1}^M \hat{y}_m I(x_i \in R_m)$$

- For regression model,  $\hat{y}_m$  can be simply the mean response in region  $R_m$ , or the regression estimate using the subset in  $R_m$ .
- For classification models, we use the sample proportions,  $p_{mc}$ , within region  $R_m$  as the estimate for the probability of class  $c$ .
- We usually assign a case to the class with the highest probability.

- The most commonly used decision tree method is Classification and Regression Tree (CART).
- To build a CART model, we first *grow* a tree using recursive *binary* splits.
- We usually stop the procedure when some stopping criterion is met. For example: the leaves must have at least  $m$  observations.
- The resulting model is typically too complex.
- Next, we *prune* the tree to obtain a simpler model and to avoid overfitting.

# Growing a Tree

- To grow a tree,
  - we choose a *cost* function, which typically reflects “impurity”
  - at each node (starting with the whole data at the root), we find the best input variable (feature),  $j^*$ , to split the data,
  - and find the best cutoff,  $t^*$  for the split.
- For numerical variables, the best  $(j^*, t^*)$  is defined as follows:

$$(j^*, t^*) = \arg \min_{j \in \{1, \dots, p\}} \min_t \{ \text{cost}(x_{ij}, y_i : x_{ij} \leq t) \\ + \text{cost}(x_{ij}, y_i : x_{ij} > t) \}$$

- When  $x_j$  is categorical with  $K$  categories, the splits are usually based on one group versus all other groups:  $x_{ij} = k$  vs.  $x_{ij} \neq k$ .

# Cost Functions

- For regression trees, the square error cost function is commonly used,

$$\text{cost}(\mathcal{D}) = \sum_{i \in \mathcal{D}} (y_i - \hat{y}_{\mathcal{D}})^2$$

- $\hat{y}_{\mathcal{D}}$  is the estimate of the response variable (e.g., mean or regression estimate) using the observations in  $\mathcal{D}$ .

# Cost Functions

- For classification trees, commonly used cost functions are
  - misclassification rate

$$\text{cost}(\mathcal{D}) = \frac{1}{|\mathcal{D}|} \sum_{i \in \mathcal{D}} (y_i \neq \hat{y}_{\mathcal{D}})$$

- Entropy

$$\text{cost}(\mathcal{D}) = \sum_{\mathcal{D}; c=1}^C p_c \log \frac{1}{p_c}$$

- Gini index

$$\text{cost}(\mathcal{D}) = \sum_{\mathcal{D}; c=1}^C p_c (1 - p_c)$$

- $p_c$  is the sample proportion of class  $c$  in the subset  $\mathcal{D}$ .

# Cost Functions




Figure: Figure 9.3 in Hastie et al. (2010) for binary classification.



# Pruning

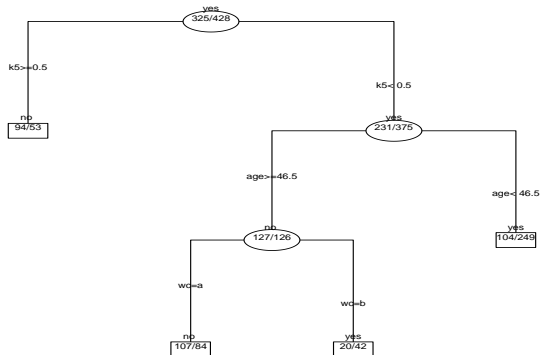
- After growing a full tree, we prune it back to avoid overfitting.
- Pruning involves collapsing some internal nodes to find a subtree without increasing the overall cost substantially.
- We can find the cross-validation costs for all possible subtrees,  $T$ , and choose the subtree with the lowest *cost-complexity* value defined as

$$C_{\alpha}(T) = \sum_{m=1}^{|T|} n_m C_m + \alpha |T|$$

where  $|T|$  is the number of terminal nodes for the subtree;  $n_m$  and  $C_m$  are the number of observations and cost for the  $m$ th terminal node;  $\alpha > 0$  is a tuning parameter.

# US Women's labor-force participation

- The dataset `Mroz` in the package `car` includes the work status of 753 women along with 7 other variables
- We obtain the following model for predicting US women's work status after growing and pruning a tree model



# Pros and Cons of Trees

- Pros
  - Easy to interpret
  - Perform automatic variable selection
  - Automatically captures interactions and nonlinear relationships
  - Robust to outliers
- Cons
  - Low predictive power
  - High variability

# Random Forests

- *Random forests* models (Breiman, 2001) attempt to reduce variance and improve predictive power by using *bagged* and *de-correlated* (decoupled) trees
- A random forests model is an ensemble of trees, each developed based on a random subset of input variables (features) and a [bootstrap] sample of observations (i.e., randomly selected observations with replacement).
- For regression, we average over the estimates from individual trees.
- For classification, each tree casts a vote, and we choose the class with the highest vote.

# Nearest Neighbor Methods

# Background

- Given a training set  $(x, y)$  of size  $n$ , we want to predict the response value,  $\tilde{y}$ , of a test case (i.e., a future observation) with input values  $\tilde{x}$
- So far, we have discussed models that build a map between the input and response variable using some parameters,  $\theta$ , after which we can forget the training set and use the resulting map to estimate  $\tilde{y}$  given  $\tilde{x}$
- A possible issue with these method is that they rely on strong assumptions (e.g., linearity, normality), which could be unrealistic
- Alternatively, we can avoid making strong assumptions and build memory-based models that remember the original training set and use it for predicting  $\tilde{y}$  directly
- To this end, we can find training cases that are close to the test case in the input space and use their average  $y$  (or median, or mode) as our estimate of  $\tilde{y}$

# $K$ -nearest neighbor

- To achieve this, we need a metric to measure closeness and specify  $k$ , the number of observations with the closest distance to the test case
- We commonly use Euclidean distance to measure closeness
- For each test case, after measuring its distance to all training cases and identifying the neighborhood  $N_k$  with the top  $k$  observations (with smallest distance to the test case), we estimate its response value as follows:

$$\hat{y}(\tilde{x}) = \frac{1}{k} \sum_{i \in N_k(x)} y_i$$

- We can use the same approach for regression, with a numerical response variable, and classification, where  $y$  is an indicator variable

# A binary classification problem with $k = 15$

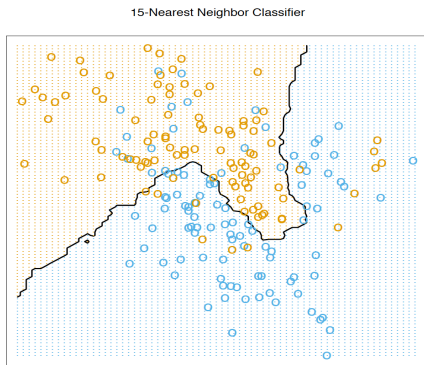


Figure: Figure 2.2 Hastie et. al. (2010).



# A binary classification problem with $k = 1$




Figure: Figure 2.3 Hastie et. al. (2010).

# Setting $k$

- As we can see in this example, the results could be very sensitive to the choice of the tuning parameter  $k$
- With small values of  $k$ , we run the risk of overfitting; with large values of  $k$ , we average over cases far from the test case
- As usual, we can use cross-validation or data splitting strategy to set this tuning parameter
- Note that although there seems to be a single parameter  $k$ , the *effective* number of parameters is  $n/k$ , i.e., the number of means we have to estimate assuming the neighborhoods are not overlapping
- Finally, note that this approach assigns 0-1 weights to the training cases; in future, we will discuss *kernel* methods where the weights are a function of distance and go smoothly to zero as distance increases

# Gaussian Process Models

- In this lecture, we discuss Gaussian process for regression.
- To learn more about this topic, refer to “Regression and classification using Gaussian process priors” (with discussion), by Neal, R. M. (1998).
- Gaussian process can be used as a distribution over functions  $y = f(x)$ .
- Note that this is a stochastic function, i.e., it includes a noise term, so even if  $x_1 = x_2$ ,  $f(x_1)$  may not be same as  $f(x_2)$  in general.

# Gaussian process models

- To introduce this concept, we start with a simple linear regression model.
- Recall that we presented a linear regression model as

$$y_i = \beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \dots \beta_p x_{ip} + \epsilon_i$$

- Using normal priors (with mean zero, and in general, different variances) for  $\beta$ 's

$$\beta_j | \sigma_j \sim N(0, \sigma_j^2) \quad j = 0, \dots, p$$

# Gaussian process models

- In prior,  $\beta$  has a  $(p + 1)$  dimensional multivariate normal distribution

$$\beta | \Sigma_\beta \sim N(0, \Sigma_\beta)$$

- $\epsilon$  also has an  $n$  dimensional multivariate normal distribution

$$\epsilon | \Sigma_\epsilon \sim N(0, \Sigma_\epsilon)$$

- To obtain the distribution of  $y$  we multiply  $\beta$  by the matrix  $x$  and add  $\epsilon$  to it.
- Based on the properties of multivariate normal distribution, the resulting distribution would still be multivariate normal  $N(0, C)$  where

$$C = x \Sigma_\beta x^T + \Sigma_\epsilon$$

# Gaussian process models

- This gives us the prior distribution on the function  $y(x)$ .
- Since any finite subset of  $y(x)$  (e.g., for the  $n$  observed cases) would have a Gaussian distribution, the prior distribution on  $y(x)$  is a *Gaussian process*.
- Similar to the Gaussian distribution, the Gaussian process is also defined by its mean (here, the mean is 0 in prior) and its covariance function  $C$ .
- For the above linear model, the elements of  $C$  are

$$C_{ij} = \text{Cov}(y_i, y_j) = \sigma_0^2 + \sum_{u=1}^p x_{iu}x_{ju}\sigma_u^2 + \delta_{ij}\sigma_\epsilon^2$$

where  $\delta_{ij}$  is equal to 1 if  $i = j$ , and 0 otherwise.

# Gaussian process models

- Setting up the model this way, we are putting the prior directly on the relationship between  $x$  and  $y$  as opposed to on some parameters that represent this relationship (i.e., we cut out the middleman).
- This is specially useful if our objective is to predict future cases as opposed to making inference about the relationship between  $x$  and  $y$ .
- Note that the prior here is implicit and reflects our choice of the functional form.
- In the above example, we are assuming the relationship is linear. In general, we could use other covariance functions,  $C$ , to create nonlinear relationship.



# Gaussian process for nonlinear regression


- For example, the following covariance function is very useful and includes a wide range of smooth nonlinear functions:

$$\text{Cov}(y_i, y_j) = \lambda^2 + \eta^2 \exp \left( - \sum_{u=1}^p \rho_u^2 (x_{iu} - x_{ju})^2 \right) + \delta_{ij} \sigma_\epsilon^2$$

- The constant part is used to make sure the model fit functions where the mean of  $y$  is not zero (the  $x$  matrix does not have a vector of 1's anymore). However, it is better to center  $y$  before analysis so we don't have to use a large constant.
- There is one  $\rho$  for each predictor.
- The noise parameter,  $\sigma_\epsilon^2$  (also called *jitter*), is essential to improve the computation.
- Within the Bayesian framework, we usually put hyper-priors on the hyper-parameters  $\lambda$ ,  $\eta$ ,  $\rho$ , and  $\sigma$ .

# The effect of parameters in the covariance function

- By using different  $\eta$ ,  $\rho$ 's,  $\lambda$  and  $\sigma_\epsilon$ , we can generate a large variety of functions.



**Figure:** The top panel shows samples based on  $\eta = 1$ ,  $\rho = 1$ ,  $\lambda = 1$ , and  $\sigma_\epsilon = 0.01$ . The bottom panel is based on the same priors except we set  $\rho = 2$ .

- As mentioned above, using a Gaussian process prior is especially useful if our goal is predicting future cases for which we only know the value of predictors,  $\tilde{x}$ .
- Assume that we have observed  $(x, y)$  for  $n$  cases, and we want to predict  $\tilde{y}$  for a new observation with predictor values  $\tilde{x}$ .

- Since the covariance function depends on  $x$ , we can find  $C_{n+1}$  for  $n$  the training cases and the new observation, i.e., for  $\begin{pmatrix} x \\ \tilde{x} \end{pmatrix}$ . To avoid confusion we denote the covariance matrix for just the training cases as  $C_n$ .
- We can write down  $C_{n+1}$  as follows:

$$C_{n+1} = \begin{pmatrix} C_n & K \\ K^T & v \end{pmatrix}$$

where  $K$  is the  $n \times 1$  covariance vector between  $\tilde{y}$  and the  $n$  observed  $y$ .  $v$  is the prior variance of  $\tilde{y}$  obtained based on the covariance function  $C$ .


- Based the above setting, we can obtain the posterior predictive distribution for the new case.
- This distribution is also Gaussian with the following mean and variance:

$$\begin{aligned}E(\tilde{y}|y) &= K^T C_n^{-1} y \\ \text{Var}(\tilde{y}|y) &= v - K^T C_n^{-1} K\end{aligned}$$

- If we need a point estimate, we can use  $E(\tilde{y}|y)$ .

# Example

- The following example shows a Gaussian process model trained on 100 data points uniformly sampled from -2 to 2 .



# Example

- For the above model, we used the following covariance function:


$$\text{Cov}(y_i, y_j) = 2 + \exp(-0.5(x_i - x_j)^2) + \delta_{ij} \times 0.1$$

- The solid line is expected function based on a grid test points between -3 and 3.
- The dashed lines show the 95% interval for predictions.

# Neural Networks



# Neural Networks (NN)



**Figure:** Multilayer perceptron (MLP) with  $p$  input variables, one hidden layer with  $q$  hidden units, and a single output. Here,  $w^{(1)}$  represents the connection weight matrix between the input layer and the hidden layer, and  $w^{(2)}$  is the vector of connection weights between the hidden layer and the output.

# Neural Networks (NN)

- A multilayer perceptron (MLP, aka feedforward NN) is comprised of an input layer, output layer and a number of hidden layers in between
- The hidden layers creates a set of basis by applying nonlinear transformations,  $g$ , to their input and pass their results to the next layer until we reach the output layer.
- We refer to  $g$  as the activation or transfer function, which is usually set to the sigmoid (aka logistic) function

$$\text{sigm}(a) = \frac{1}{1 + e^{-a}}$$

or the hyperbolic tangent function

$$\tanh(a) = \frac{e^a - e^{-a}}{e^a + e^{-a}}$$

# Neural Networks (NN)




Figure: Figure 16.6 in Murphy (2012).

# Neural Networks (NN)

- The output  $f(x)$ , which is a function used to approximate  $y$ , is a linear combination of basis defined by the hidden layers
- For linear regression models (continuous outcome)

$$P(y|x, w) = N(y|f(x), \sigma^2)$$

- For logistic regression models (binary outcome)

$$P(y|x, w) = \text{Ber}(y|\text{sigm}(f(x)))$$

- For multiple categories, we use the multinomial logit model, which is also known as the softmax function

# Neural Networks (NN)

- For a MLP with one hidden layer and tanh activation function, we have

$$g_j(x) = \tanh[w_{0j}^{(1)} + \sum_{k=1}^p w_{kj}^{(1)} x_k], \quad \text{for } j = 1, \dots, q$$

$$f(x) = w_0^{(2)} + \sum_{j=1}^q w_j^{(2)} g_j(x)$$

- Here  $w_0$ 's, which play the role of the intercept in regression models, are called biases

- To train a MLP, we first need to specify the negative log-likelihood, which is also known as the energy function,  $E$
- For regression, we have (squared error)

$$E = \frac{1}{2\sigma^2} \sum_{i=1}^n (y_i - f(x_i))^2$$

for binary classification models, we have (cross entropy)

$$E = - \sum_{i=1}^n y_i [\log \text{sigm}(f(x_i))] + (1 - y_i) [\log(1 - \text{sigm}(f(x_i)))]$$

- To estimate the weights, we minimize the energy function with respect to  $w$

- The parameters of a neural network model are not identifiable
  - Permuting the order of hidden units does not change the model
  - If we change the sign of weights entering a hidden unit, the model remains the same as long as we also change the sign of the weights going out of that unit since  $\tanh(-a) = -\tanh(a)$
- Also, in general the energy function for MLP is non-convex
- Nevertheless, we can still use common iterative optimization methods (e.g., gradient descent algorithms) to obtain locally optimal estimates
- Using the chain rule, it is easy to find the gradient

# Backpropagation

- For learning (i.e., estimating the weights), we start with initializing the weights (including the biases) to some random numbers (all different values) and iteratively perform the following steps
  - At each iteration, we use *forward propagation* to find the values going to each unit, before and after transformation, until we reach the output layer
  - We find the derivatives of  $E$  with respect to each unit starting from the output and *backpropagate* using the chain rule to find the derivatives with respect to hidden units
  - Using the chain rule again, we find the derivatives with respect to the weights
  - We then update the weights by moving in the direction of the negative gradient (see the notes on optimization)
  - We repeat the above steps until some stopping criterion is reached



# Backpropagation

- For a MLP with one hidden layer and tanh activation function, given the current weights, forward propagation involves finding the following values at each hidden unit before and after transformation

$$z_j = w_{0j}^{(1)} + \sum_{k=1}^p w_{kj}^{(1)} x_k$$
$$g_j = \tanh(z_j)$$

- For the output unit, we have

$$f = w_0^{(2)} + \sum_{j=1}^q w_j^{(2)} g_j$$

# Backpropagation

- Backpropagation starts with finding  $\partial E / \partial f$ ; For regression model

$$\frac{\partial E}{\partial f} = -\frac{1}{\sigma^2} \sum_{i=1}^n (y - f(x_i))$$

- Next, using the chain rule we have

$$\frac{\partial E}{\partial g_j} = \frac{\partial E}{\partial f} \frac{\partial f}{\partial g_j} = w_j^{(2)} \frac{\partial E}{\partial f}$$

- We then find the derivatives with respect to  $z_j$

$$\frac{\partial E}{\partial z_j} = \frac{\partial E}{\partial g_j} \frac{\partial g_j}{\partial z_j} = (1 - g_j^2) \frac{\partial E}{\partial g_j}$$

Recall that  $\frac{d}{da} \tanh(a) = 1 - \tanh^2(a)$

# Backpropagation

- Finally, we find the derivatives with respect to the weights
- For connection weights between the hidden layer and output we have

$$\frac{\partial E}{\partial w_j^{(2)}} = \frac{\partial E}{\partial f} \frac{\partial f}{\partial w_j^{(2)}} = g_j \frac{\partial E}{\partial f}$$

Note that  $g_0 = 1$  when evaluating  $\partial E / \partial w_0^{(2)}$

- For the weights connecting the input layer to the hidden layer we have

$$\frac{\partial E}{\partial w_{kj}^{(1)}} = \frac{\partial E}{\partial z_j} \frac{\partial z_j}{\partial w_{kj}^{(1)}} = x_k \frac{\partial E}{\partial z_j}$$

$x_0 = 1$  when evaluating  $\partial E / \partial w_0^{(1)}$

# Backpropagation

- After we find the gradient  $\nabla E$ , we update the parameters by taking a step in a direction of negative gradient, with stepsize  $t$

$$\begin{aligned}\Delta w &= -\nabla E \\ w &\leftarrow w + t\Delta w\end{aligned}$$

- The stepsize (aka learning rate) is found by trial-and-error
- We could run the algorithm until the approximation error falls below a desired threshold; however, this could lead to *overfitting*
- Two common strategies to avoid this issue are *early stopping* and *weight decay*

# Early stopping

- Overfitting occurs when the model performs well on the training data and performs poorly on the test (future) data
- In the “early stopping” method, we start with some initial weights close to zero and monitor the performance of the neural network model throughout the training process based on an independent *validation set* (usually 20% of the data; this is separate from any test set used for model evaluation); we stop the algorithm when the model's performance on the validation set starts to decline substantially (a sign of overfitting)
- We can use the prediction error or average log probability on the validation set as a measure of performance
- This method could be very successful for avoiding overfitting; however it is *ad hoc* and wasteful since some of the data points are not used in the training directly

- Alternatively, to avoid overfitting, we can penalize models against complexity (similar to ridge regression and Lasso)
- To this end, instead of minimizing the energy function, we minimize the penalized version of it by adding the following penalty terms:

$$\lambda_1 \sum_{k=1}^p \sum_{j=1}^q [w_{kj}^{(1)}]^2 + \lambda_2 \sum_{j=1}^q [w_j^{(2)}]^2$$

- This is known as “weight decay” since it shrinks the weights towards zero to encourage simpler models
- To set the values of  $\lambda_1$  and  $\lambda_2$ , we can use an independent validation set as before or use cross-validation when the sample size is small

# Illustrative Example




Figure: Observed data

# Illustrative Example




Figure: A neural network with 5 hidden units



# Illustrative Example




Figure: Estimated function,  $f(x)$ , with 5 hidden units

# Illustrative Example




Figure: Estimated function,  $f(x)$ , with 50 hidden units

# Illustrative Example




Figure: Estimated function,  $f(x)$ , with 50 hidden units using weight decay